

# Task 1: Frequency Analysis

COMP3028 Computer Security — Coursework 1

David Leong Jin Tao  
Student ID: 20620891

February 25, 2026

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Task Statement . . . . .	3
1.2	Scope and Structure . . . . .	3
<b>2</b>	<b>Literature Review</b>	<b>3</b>
2.1	Role of n-gram statistics . . . . .	3
2.2	Fitness (scoring) functions . . . . .	3
2.3	Search algorithms . . . . .	4
2.4	Effectiveness and limitations . . . . .	4
<b>3</b>	<b>Methodology</b>	<b>4</b>
3.1	Design choices . . . . .	5
<b>4</b>	<b>Implementation and code</b>	<b>5</b>
4.1	Obtaining reference data from Wikipedia . . . . .	5
4.2	Loading letter frequency data . . . . .	6
4.3	Decryption and key application . . . . .	7
4.4	Letters-only and n-gram extraction . . . . .	7
4.5	Cipher letter counts . . . . .	7
4.6	Building the initial key from letter frequency . . . . .	8
4.7	Crib-based initial key . . . . .	8
4.8	Unigram probabilities and safe log . . . . .	8

4.9	Fitness function . . . . .	9
4.10	Hill climbing . . . . .	9
4.11	Multiple restarts and best-key selection . . . . .	10
4.12	Exhaustive pair-swap refinement . . . . .	10
4.13	Manual swap application . . . . .	11
4.14	Output format . . . . .	11
<b>5</b>	<b>Results</b>	<b>12</b>
5.1	Recovered encryption key . . . . .	12
5.2	Evidence of correct decryption . . . . .	12
5.3	Quantitative results . . . . .	13
<b>6</b>	<b>Program output and execution log</b>	<b>13</b>
6.1	Heuristic phase: from loading data to exhaustive refinement . . . . .	13
6.2	Manual phase: letter-swap corrections and final key . . . . .	15
<b>7</b>	<b>Observations and discussion</b>	<b>16</b>
7.1	Observations . . . . .	16
7.2	Comparison with the SEED Lab exercise . . . . .	17
<b>8</b>	<b>Conclusion</b>	<b>17</b>

# 1 Introduction

## 1.1 Task Statement

This report addresses Task 1 of COMP3028 Coursework 1: to decipher the substitution ciphertext provided on Moodle and to recover the encryption key. The ciphertext was produced by a *monoalphabetic substitution cipher*, in which each plaintext letter is replaced by a unique ciphertext letter according to a fixed permutation of the 26-letter alphabet.

## 1.2 Scope and Structure

The report presents the approach adopted, the implementation, and the results. It is organised as follows: (i) the use of monogram, bigram, and trigram frequency data to construct and score candidate keys; (ii) the methodology, including the initial key from frequency ranking and an optional crib, hill climbing with multiple restarts, exhaustive pair-swap refinement, and limited manual correction; (iii) the implementation, with relevant code and justification; (iv) the recovered encryption key and decrypted plaintext, with supporting program output; (v) observations and a comparison with the SEED Lab exercise; and (vi) conclusions.

# 2 Literature Review

Established methods for breaking monoalphabetic substitution ciphers rely on *n*-gram frequency analysis: monograms (single letters), bigrams, trigrams, and optionally quadgrams [1, 2].

## 2.1 Role of n-gram statistics

In English, letter frequencies follow a characteristic distribution (e.g. E, T, A, O, I, N, S, H, R, ...). Mapping the most frequent cipher letter to E, the second to T, and so on yields a first approximation of the key. Bigrams (e.g. TH, HE, IN, ER) and trigrams (e.g. THE, AND, ING) exhibit more skewed distributions; correct plaintext contains many high-probability n-grams, whereas incorrect keys produce rare n-grams [1]. Quadgrams further improve discrimination at the cost of additional data and computation [2]. Reference statistics are typically drawn from corpora such as Norvig's letter n-gram counts [3].

## 2.2 Fitness (scoring) functions

A fitness function assigns a *higher* score to text that is more English-like. Probabilities (or log-probabilities) for 1-grams, 2-grams, and 3-grams are usually precomputed from a reference corpus. For a candidate decryption, each n-gram is looked up; unknown or zero-probability n-grams are assigned a small floor probability to avoid  $\log(0)$ . The score is typically the **sum of log-probabilities** (log-likelihood) over monograms, bigrams, and trigrams, optionally combined as a weighted sum [1, 4]. The objective is to find the key that *maximises* this fitness when applied to the ciphertext.

## 2.3 Search algorithms

The key space has size  $26!$ , so exhaustive search is infeasible. **Hill climbing** starts from an initial key (e.g. from frequency ranking or random) and repeatedly moves to a *neighbour* key that improves fitness; a natural neighbourhood is **swapping two letters** in the key ( $\binom{26}{2} = 325$  neighbours). The process terminates at a local maximum [1]. Because hill climbing may terminate at a local optimum, **multiple restarts** from different initial keys are used, and the best key over all runs is retained [2, 1]. Some variants use targeted neighbours (e.g. swapping letters to correct the least common bigrams in the current decryption) or **simulated annealing** to accept worse keys with decreasing probability so as to escape local optima.

## 2.4 Effectiveness and limitations

The approach is effective for monoalphabetic substitution given sufficiently long ciphertext; the *unicity distance* is approximately 28 characters [2]. Short or statistically atypical plaintext may resist recovery or remain ambiguous. Success depends on the reference corpus matching the plaintext language.

Building on these ideas, the following section describes the methodology adopted for the present task.

# 3 Methodology

The methodology comprised the following steps.

1. **Input.** The ciphertext was read from `data/task1/article_encrypted.txt` (1530 characters in total). Non-letter characters (e.g. [] and newlines) were preserved unchanged; only letters a–z were subject to the substitution mapping.
2. **Reference data.** Letter (monogram), bigram, and trigram frequencies were loaded from `data/frequencies/*.json`. These JSON files can be produced by running `scripts/fetch_wikipedia_frequencies.py`, which fetches the corresponding frequency tables from Wikipedia (Letter frequency, Bigram, Trigram) and saves them in the format expected by the decrypt script. Letter frequencies defined the plaintext letter order (E, T, A, O, …); bigram and trigram data were converted to probabilities for the fitness function. Unknown n-grams were assigned small default probabilities ( $10^{-5}$ ,  $10^{-8}$ ,  $10^{-10}$  for unigrams, bigrams, and trigrams respectively) so as to avoid  $-\infty$  in the log-sum.
3. **Initial key.** Ciphertext letter counts were computed and ranked by frequency. One initial key was constructed by mapping the most frequent cipher letter to E, the second to T, and so on. An alternative initial key was constructed using a crib when the most frequent cipher trigram (e.g. `meh`) had count at least 3: that trigram was assumed to correspond to `the`, and the remaining key positions were filled by frequency rank.
4. **Fitness function.** For each candidate key, the ciphertext was decrypted and the letters-only string was scored. Fitness was defined as a weighted sum of log-probabilities of unigrams, bigrams, and trigrams (weights 0.2, 0.3, 0.5). Higher fitness corresponds to more English-like text.

5. **Hill climbing.** From each initial key, hill climbing was applied: two randomly chosen positions in the key were swapped, the ciphertext was decrypted, and fitness was recomputed; the swap was retained only if fitness improved. The process ran for up to 10 000 iterations and terminated after 1500 iterations without improvement. Five restarts were run, alternating between the frequency-based and crib-based initial keys; the key with the highest fitness over all runs was retained.
6. **Exhaustive refinement.** The best key was refined by enumerating all 325 pair swaps; any swap that improved fitness was retained, and the process was repeated until no further improvement was possible.
7. **Final key.** A small set of manual letter swaps was applied to correct remaining confusions (e.g. b/w, k/v), yielding a fully coherent decryption. The final key was written in the required format and used to produce the decrypted plaintext file.

### 3.1 Design choices

- **Trigrams and bigrams.** Trigrams were used because they discriminate strongly between correct and incorrect decryptions; the implementation also used bigrams and monograms in a single weighted fitness to stabilise the search.
- **Multiple restarts.** Hill climbing may converge to a local maximum; multiple restarts from different seeds and from both frequency-based and crib-based keys increase the probability of recovering the correct key.
- **Unknown n-grams.** Assigning a small floor probability avoids  $\log(0)$  and ensures that the fitness is defined for any decryption; rare or missing n-grams then contribute a fixed negative penalty.

## 4 Implementation and code

The methodology above was implemented in a single Python script, `scripts/decrypt_task1.py`. The frequency JSON files used by the script are produced by a separate script, `scripts/fetch_wikipedia_frequencies.py`, which obtains letter, bigram, and trigram data from Wikipedia and saves them as JSON. The following subsections describe first how that reference data is obtained and saved, then how the decrypt script uses it: loading letter frequency data; decryption and key application; letter-only and n-gram extraction; cipher letter counts and initial key construction; unigram probabilities and safe logarithm; the fitness function; hill climbing; multiple restarts and best-key selection; exhaustive pair-swap refinement; manual swap application; and the output format.

### 4.1 Obtaining reference data from Wikipedia

(`scripts/fetch_wikipedia_frequencies.py`, `fetch_page` lines 29–33, `main` lines 122–147.) The letter, bigram, and trigram frequency data are fetched from the following Wikipedia pages and saved as JSON in `data/frequencies/`:

- **Letter frequency:** [https://en.wikipedia.org/wiki/Letter\\_frequency](https://en.wikipedia.org/wiki/Letter_frequency) — the main English letter frequency table (Letter, Relative frequency in texts, in dictionaries) is

parsed from a `wikitable`; each row yields `letter`, `texts_percent`, and optionally `dictionaries_percent`.

- **Bigram:** <https://en.wikipedia.org/wiki/Bigram> — bigram percentages are taken from a `<pre>` block on the page (e.g. th 3.56%, he 3.07%); each pair is stored as `bigram` and `frequency_percent`.
- **Trigram:** <https://en.wikipedia.org/wiki/Trigram> — the trigram frequency table (Rank, Trigram, Frequency) is parsed from an HTML table; each row yields `trigram` and `frequency_percent`.

The script uses `requests` to fetch each page (with a User-Agent header) and `BeautifulSoup` to parse the HTML. Parsed data are written with `json.dump` to `data/frequencies/letter_frequency.json`, `data/frequencies/bigram_frequency.json`, and `data/frequencies/trigram_frequency.json`. The decrypt script then reads these files and expects the field names above (`texts_percent` for letters, `frequency_percent` for bigrams and trigrams).

Listing 1: Fetching a page and saving letter frequency as JSON.

```

1 def fetch_page(url: str) -> str:
2     resp = requests.get(url, headers={"User-Agent": USER_AGENT}, timeout
3                         =30)
4     resp.raise_for_status()
5     return resp.text
6
7 # In main(): for each of letter, bigram, trigram:
8 #     html = fetch_page(URL_...)
9 #     data = parse_letter_frequency(html) # or parse_bigram_*, 
10 #          parse_trigram_*
11 #     with open(OUTPUT_DIR / "letter_frequency.json", "w", encoding="utf
12 #             -8") as f:
13 #                 json.dump(data, f, indent=2)

```

## 4.2 Loading letter frequency data

(scripts/decrypt\_task1.py, `load_letter_frequencies`, lines 61–87.) The script reads `letter_frequency.json`, parses each letter and its `texts_percent` value, and sorts by frequency descending to obtain the plaintext letter order (E, T, A, O, ...). Invalid percentage values are caught and set to 0. Bigram and trigram data are loaded similarly from their JSON files (using `frequency_percent`) and converted to probability dictionaries for the fitness function.

Listing 2: Loading letter frequency and ordering by frequency.

```

1 with open(LETTER_FREQ_PATH, encoding="utf-8") as f:
2     data = json.load(f)
3 out = []
4 for item in data:
5     letter = item["letter"].upper()
6     raw = item.get("texts_percent", "0%").replace("%", "").strip()
7     try:
8         pct = float(raw)
9     except ValueError:
10        pct = 0.0
11    out.append((letter.lower(), pct))
12 out.sort(key=lambda x: -x[1]) # E, T, A, 0, ...

```

### 4.3 Decryption and key application

(scripts/decrypt\_task1.py, decrypt\_text, lines 167–171.) The key is represented as a 26-character string where `key[i]` is the cipher letter for plain letter `ALPHABET[i]` (i.e. encryption: plain `a` → `key[0]`). Decryption uses the inverse mapping: for each cipher letter `c`, find the index `i` such that `key[i]==c`, then the plain letter is `ALPHABET[i]`.

Listing 3: Decryption (inverse substitution).

```
1 def decrypt_text(ciphertext: str, key: str) -> str:
2     inv = {key[i]: ALPHABET[i] for i in range(26)}
3     return "".join(inv.get(c.lower(), c) if c.isalpha() else c for c in
4                     ciphertext)
```

Every candidate key is evaluated by decrypting the ciphertext and scoring the result; this mapping is therefore central to the key-recovery process.

### 4.4 Letters-only and n-gram extraction

(scripts/decrypt\_task1.py, lines 151–159.) The fitness function scores only letters; non-letters (e.g. `[]`, newlines) are stripped before counting. Overlapping n-grams are extracted with a sliding window so that every bigram and trigram in the decrypted text contributes to the score.

Listing 4: Letter-only text and n-gram extraction.

```
1 def letters_only(text: str) -> str:
2     return "".join(c.lower() for c in text if c.isalpha())
3
4 def extract_ngrams(text: str, n: int) -> list[str]:
5     t = letters_only(text)
6     return [t[i : i + n] for i in range(len(t) - n + 1)] if len(t) >= n
7     else []
```

**Rationale.** Cipher letter frequency ranking and the crib (most common cipher trigram) both operate on letter-only input; the fitness function uses `extract_ngrams` to obtain bigrams and trigrams from the candidate decryption. Omitting this step would allow non-letter characters to distort the statistics.

### 4.5 Cipher letter counts

(scripts/decrypt\_task1.py, cipher\_letter\_counts, lines 182–188.) The ciphertext (letters only) is counted per letter and sorted by count descending. This ranking is later matched to English letter frequency order to build the initial key.

Listing 5: Rank cipher letters by frequency.

```
1 def cipher_letter_counts(letter_only: str) -> list[tuple[str, int]]:
2     cnt = Counter(letter_only)
3     return sorted(
4         [(c, cnt[c]) for c in ALPHABET if cnt[c] > 0],
5         key=lambda x: -x[1],
6     )
```

**Rationale.** This function produces the cipher-side frequency ranking that `build_initial_key` maps to the plaintext letter order (E, T, A, ...), under the assumption that the most frequent cipher letter corresponds to E, the second to T, and so on.

## 4.6 Building the initial key from letter frequency

(`scripts/decrypt_task1.py`, `build_initial_key`, lines 191–223.) Cipher letters are ranked by count; plain letters are ordered by reference frequency (E, T, A, ...). The most frequent cipher letter is mapped to E, the second to T, and so on.

Listing 6: Initial key from frequency ranking (excerpt).

```

1 plain_ranked = [x[0] for x in letter_freq_order]
2 cipher_letters_ranked = [x[0] for x in cipher_ranked]
3 # ... build key_list so key_list[ALPHABET.index(plain)] = cipher
4 return "".join(key_list)

```

This provides a principled starting point for hill climbing, avoiding the need to start from an arbitrary permutation.

## 4.7 Crib-based initial key

(`scripts/decrypt_task1.py`, `build_initial_key_with_crib`, lines 226–270.) When the most frequent cipher trigram (e.g. `meh`) is assumed to be `the`, the three letter positions are fixed in the key; the remaining 23 positions are filled by matching plain letter frequency order to the remaining cipher letters by rank.

Listing 7: Crib: fix three positions, fill rest by frequency.

```

1 # Fix crib: e.g. cipher "meh" -> plain "the"
2 for i in range(3):
3     plain = plain_trigram[i]
4     cipher = cipher_trigram[i]
5     idx = ALPHABET.index(plain)
6     key_list[idx] = cipher
7 # Fill remaining positions from plain_ranked and cipher_letters_ranked
8 # (skip already fixed), then any gaps from unused alphabet.

```

**Rationale.** The crib supplies an alternative initial key that typically attains a better local maximum than the purely frequency-based key; the main loop executes hill climbing from both and retains the key with the highest fitness.

## 4.8 Unigram probabilities and safe log

(`scripts/decrypt_task1.py`, `safe_log`, lines 276–278; `unigram_probs_from_letter_freq`, lines 431–435.) Reference letter frequencies are normalised to probabilities (sum to 1) for the fitness. To avoid  $\log(0)$  when an n-gram is missing from the reference, a floor value is used.

Listing 8: Unigram probabilities from reference; safe logarithm.

```

1 def safe_log(p: float, floor: float = 1e-10) -> float:
2     return math.log(max(p, floor))
3

```

```

4 def unigram_probs_from_letter_freq(letter_freq) -> dict[str, float]:
5     total = sum(p for _, p in letter_freq)
6     if total <= 0:
7         total = 1.0
8     return {letter: p / total for letter, p in letter_freq}

```

**Rationale.** The fitness function uses these unigram probabilities and `safe_log` for every letter and n-gram; without the floor, unknown n-grams would yield  $-\infty$  and invalidate the comparison between candidate keys.

## 4.9 Fitness function

(`scripts/decrypt_task1.py`, `fitness`, lines 281–305.) The fitness is a weighted sum of log-probabilities for unigrams, bigrams, and trigrams in the letter-only decrypted text. Missing n-grams use default probabilities (and `safe_log` avoids  $-\infty$ ).

Listing 9: Fitness (n-gram log-probability sum).

```

1 def fitness(decrypted_letters, unigram_probs, bigram_probs,
2             trigram_probs,
3             weights=(0.2, 0.3, 0.5)) -> float:
4     default_uni = 1e-5
5     default_bi = 1e-8
6     default_tri = 1e-10
7     score = 0.0
8     for c in decrypted_letters:
9         score += weights[0] * safe_log(unigram_probs.get(c, default_uni))
10    for bg in extract_ngrams(decrypted_letters, 2):
11        score += weights[1] * safe_log(bigram_probs.get(bg, default_bi))
12    for tg in extract_ngrams(decrypted_letters, 3):
13        score += weights[2] * safe_log(trigram_probs.get(tg, default_tri))
14
15    return score

```

Maximising this score is the objective of the search and is the criterion by which candidate keys are compared.

## 4.10 Hill climbing

(`scripts/decrypt_task1.py`, `hill_climb`, lines 311–367.) The algorithm repeatedly swaps two random positions in the key; if the new key yields higher fitness, it is retained; otherwise the swap is reverted.

Listing 10: Hill climbing with random pair swaps.

```

1 i, j = random.randint(0, 25), random.randint(0, 25)
2 if i == j: continue
3 key_list[i], key_list[j] = key_list[j], key_list[i]
4 new_key = ''.join(key_list)
5 new_score = fitness(letters_only(decrypt_text(cipher_letters, new_key)),
6                     ...)
7 if new_score > current_score:
8     current_key = new_key
9     current_score = new_score

```

```

9     no_improve = 0
10    else:
11        key_list[i], key_list[j] = key_list[j], key_list[i]
12        no_improve += 1

```

This refines the key by using n-gram fitness to correct local errors that arise from monogram-only frequency mapping.

## 4.11 Multiple restarts and best-key selection

(scripts/decrypt\_task1.py, main, lines 518–547.) Hill climbing is executed five times from different starting keys (alternating frequency-based and crib-based); each run uses a fixed random seed for reproducibility. The key with the highest fitness over all runs is retained for exhaustive refinement.

Listing 11: Multiple restarts; keep best key by fitness.

```

1 candidates = [(initial_key, "frequency-based")]
2 if initial_key_crib is not None:
3     candidates.append((initial_key_crib, "crib meh->the"))
4 best_key, best_score = None, -float("inf")
5 for run in range(num_restarts):
6     start_key, _ = candidates[run % len(candidates)]
7     key = hill_climb(ciphertext_raw, start_key, ...)
8     score = fitness(letters_only(decrypt_text(ciphertext_raw, key)),
9                      ...)
10    if score > best_score:
11        best_score, best_key = score, key

```

**Rationale.** A single hill-climb may terminate at a local maximum; multiple restarts from both frequency-based and crib-based initial keys, with retention of the best result, increase the probability of recovering the correct substitution.

## 4.12 Exhaustive pair-swap refinement

(scripts/decrypt\_task1.py, exhaustive\_swap\_refinement, lines 372–425.) After hill climbing, every pair of indices  $(i, j)$  with  $i < j$  is tried in order; the first swap that improves fitness is kept, then the full enumeration is repeated. This continues until no single swap improves fitness, reaching a local maximum over all 325 possible single swaps.

Listing 12: Exhaustive refinement: try all pair swaps.

```

1 key_list = list(current_key)
2 for i in range(26):
3     for j in range(i + 1, 26):
4         key_list[i], key_list[j] = key_list[j], key_list[i]
5         new_key = ''.join(key_list)
6         new_score = fitness(letters_only(decrypt_text(ciphertext,
7                                         new_key)), ...)
8         if new_score > current_score:
9             current_key = new_key
10            current_score = new_score
11            improved = True
12            break
13        else:

```

```
13     key_list[i], key_list[j] = key_list[j], key_list[i]
```

**Rationale.** Hill climbing uses random swaps and may halt before testing every pair; exhaustive refinement guarantees that no improving single swap is overlooked, thereby correcting remaining letter confusions (e.g. b/w, p/k) that the stochastic search may have missed.

### 4.13 Manual swap application

(scripts/decrypt\_task1.py, apply\_swaps, lines 440–445.) After exhaustive refinement, a fixed list of index pairs is applied to correct remaining letter confusions (e.g. b↔w, k↔v) that the n-gram fitness did not resolve. Each pair  $(i, j)$  swaps `key[i]` and `key[j]`; the resulting key is used for the final decryption.

The swaps were applied in the following order (plain-letter indices: a=0, b=1, …, z=25), with the reason for each:

1. (1, 22): b↔w. Corrects “world” vs. “was” (e.g. “end of world war i”).
2. (10, 21): k↔v. Corrects “marketing” vs. “governments” (e.g. “military and government services”).
3. (1, 15): b↔p. Corrects “product” vs. “described” (e.g. “finished product”, “described in detail”).
4. (9, 23): j↔x. Corrects “adjacently” (e.g. “used … adjacently”).
5. (16, 23): q↔x. Corrects “complex” (e.g. “the most complex”).
6. (16, 25): q↔z. Corrects “blitzkrieg” and “emphasized”.

Listing 13: Apply a list of index swaps to the key.

```
1 def apply_swaps(key: str, swaps: list[tuple[int, int]]) -> str:
2     key_list = list(key)
3     for i, j in swaps:
4         key_list[i], key_list[j] = key_list[j], key_list[i]
5     return ''.join(key_list)
```

### 4.14 Output format

(scripts/decrypt\_task1.py, format\_key\_report, lines 448–456.) The key is written in the coursework format: one line of plaintext letters, one line of ciphertext letters (same order).

Listing 14: Key output format.

```
1 def format_key_report(key: str) -> str:
2     plain_line = "Plaintext " + ".join(ALPHABET)
3     cipher_line = "Ciphertext " + ".join(key)
4     return plain_line + "\n" + cipher_line
```

This produces the key in the format required for the report.

## 5 Results

Application of the methodology and implementation described in Sections 3 and 4 to the provided ciphertext yielded the results below. This section presents the recovered encryption key, evidence of correct decryption (in both raw and spaced form), and quantitative summary statistics. Section 6 reproduces the full execution log.

### 5.1 Recovered encryption key

The recovered encryption key maps each plaintext letter to a ciphertext letter. Table 1 gives the mapping in the format required by the coursework (plaintext a–z in order, ciphertext letters in the same order).

Plaintext	a	b	c	d	e	f	g	h	i	j	k	l	m	n	o	p	q	r	s	t	u	v	w	x	y	z
Ciphertext	g	t	q	d	h	b	j	e	i	v	z	r	c	w	s	p	f	x	u	m	k	o	l	y	n	a

Table 1: Recovered substitution key: plaintext letter → ciphertext letter (e.g. a→g, b→t, z→a).

### 5.2 Evidence of correct decryption

The decrypted output was written to `data/task1/article_decrypted.txt`. Because the ciphertext contained no spaces (see Section 7), the decryption is a continuous letter stream. A representative excerpt (first 500 characters) from that file is reproduced verbatim below:

```
theenigmamachinewasinventedbythegermanengineerarthurscherbiusatthe  
endofworldwari[]thiswasnotknownuntilwhenapaperbykarlde  
leeuwwasfoundthatdescribedindetailarthurscherbiuschanges[]the  
germanfirmscherbiusrittercofoundedbyarthurscherbiuspatentedideas  
foraciphermachineinandbegankmarketingthefinishedproductunder  
thebrandnameenigmaininitiallytargetedadcommercialmarkets[]  
earlymodelsWEREusedcommerciallyfromtheearlysandadoptedby  
militaryandgovernmentservicesofseveralcountriesmostnotablynazi  
germanybeforeandduringworldwarii  
  
severaldifferentenigmamodelswereproducedbutthegermanmilitarymodels  
havingaplugboardwerethemostcomplexjapaneseandalianmodelswere
```

The content is recognisably about the Enigma machine (Arthur Scherbius, World War I and II, commercial and military use). The [] symbols are preserved from the ciphertext as non-letter characters that were not substituted.

To demonstrate that this letter stream corresponds to readable English, the same plaintext with word boundaries restored (from `data/task1/article_decrypted_spaced.txt`) is given below. The decryption is unchanged; spacing has been added for readability only.

```
the enigma machine was invented by the german engineer arthur scherbius at the  
end of world war i [] this was not known until when a paper by karl de  
leeuw was found that described in detail arthur scherbius changes [] the  
german firm scherbius ritter co founded by arthur scherbius patented ideas
```

for a cipher machine in and began marketing the finished product under the brand name enigma in initially targeted at commercial markets [] early models were used commercially from the early s and adopted by military and government services of several countries most notably nazi germany before and during world war ii

several different enigma models were produced but the german military models having a plugboard were the most complex japanese and italian models were also in use with its adoption in slightly modified form by the german navy in and the german army and air force soon after the name enigma became widely known in military circles pre war german military planning emphasized fast mobile forces and tactics later known as blitzkrieg which depend on radio communication for command and coordination since adversaries would likely intercept radio signals messages would have to be protected with secure encoding compact and easily portable the enigma machine filled that need

### 5.3 Quantitative results

- Ciphertext length: 1530 characters.
- Cipher letter counts (top 5): h=188, g=142, w=120, i=118, x=109.
- Most common cipher trigram: `meh` (count 21), used as crib for `the`.
- Restarts: 5 (alternating frequency-based and crib-based initial keys).
- Final fitness (best run): Restarts 2 and 4 (crib-based) achieved fitness values of approximately  $-22974$  to  $-22977$ ; frequency-based restarts achieved approximately  $-23921$  to  $-24100$ .
- Exhaustive pair-swap refinement was applied after hill climbing; manual swaps were then applied to achieve full coherence.

## 6 Program output and execution log

The key and decrypted text presented in Section 5 were produced by executing `pythonscripts/decrypt_task1.py`. The output below confirms the recovered key, manual swaps, and decrypted plaintext. To satisfy the requirement for complete steps with justification, the full terminal output is reproduced below as plain text (no screenshots). The log is divided into two parts: the heuristic phase (frequency analysis, initial key construction, hill climbing, and exhaustive refinement), and the manual phase (letter-swap corrections and final key).

### 6.1 Heuristic phase: from loading data to exhaustive refinement

The execution log from script start through exhaustive pair-swap refinement is given below. Each phase includes the STEP and JUSTIFICATION output by the script.

```
1 COMP3028 Task 1: Substitution cipher decryption using frequency analysis
2 Using provided letter, bigram and trigram data with full process logging.
3
```

```

4 =====
5   Loading letter frequency data
6 =====
7   Loaded 26 letters; order by frequency (highest first): e t a o i n s h r d ...
8   STEP: Use letter frequency order to guess substitution.
9   JUSTIFICATION: In English, E is most common (~12.7%), then T (~9.1%), A, O, I,
10  N, etc. We will map the most frequent cipher letter to E, the second to T,
11  and so on.
12
13 =====
14   Loading bigram frequency data
15 =====
16   Loaded 42 bigrams; e.g. th=0.0356, he=0.0307
17   STEP: Use bigram frequencies to score and refine the key.
18   JUSTIFICATION: Common English bigrams (e.g. th, he, in, er) should appear
19      often
20      in correct plaintext; wrong keys produce rare bigrams. We score decrypted text
21      by sum of log(bigram_prob).
22
23 =====
24   Loading trigram frequency data
25 =====
26   Loaded 16 trigrams; e.g. the=0.0181, and=0.0073
27   STEP: Use trigram frequencies to score and refine the key.
28   JUSTIFICATION: Trigrams like 'the', 'and', 'ing' are very common in English;
29      incorrect decryptions contain rare trigrams and get a lower score.
30
31 =====
32   Loading ciphertext
33 =====
34   Loaded 1530 characters from article_encrypted.txt
35   STEP: We preserve the full ciphertext including non-letters (e.g. []).
36   JUSTIFICATION: Non-letters are left unchanged when applying the substitution;
37      only a-z are decrypted.
38
39 =====
40   Analysing ciphertext letter frequencies
41 =====
42   Cipher letter counts (top 10): h=188, g=142, w=120, i=118, x=109, m=107,
43      s=89, u=86, d=70, r=63
44   STEP: Count how often each cipher letter appears.
45   JUSTIFICATION: This ranking is matched to English letter frequency ranking to
46      form the initial key.
47
48 =====
49   Building initial key from letter frequency (monogram)
50 =====
51   Plain (by freq): e t a o i n s h r d l c ...
52   Cipher (by freq): h g w i x m s u d r c q ...
53   STEP: Map cipher letters to plain letters by matching frequency rank.
54   JUSTIFICATION: The most frequent cipher letter is assumed to be E, the second
55      T, etc. This gives a first approximation of the substitution key.
56
57   Most common cipher trigram: 'meh' (count=21)
58   STEP: Also build an alternative initial key using crib: 'meh' -> 'the'.
59   JUSTIFICATION: The most frequent trigram in English is 'the'; using it to
60      seed the key often improves the result.
61
62 =====
63   Refining key with hill climbing (bigram and trigram)
64 =====
65   STEP: Run several restarts from frequency-based and crib-based keys; keep the
       key with highest fitness.

```

```

66 JUSTIFICATION: Hill climbing can stop at a local maximum; multiple restarts
67 increase the chance of finding the true key.
68
69 Initial fitness (log-probability sum): -24710.75
70 STEP: Iteratively swap two letters in the key and re-score decrypted text.
71 JUSTIFICATION: If the new key yields higher n-gram fitness, we keep it. This
72 escapes local errors from monogram-only mapping (e.g. similar-frequency
73 letters).
74
75 Iteration 0: fitness improved to -24413.63 (swap key[u]<->key[d])
76 Iteration 16: fitness improved to -24338.60 (swap key[o]<->key[s])
77 Iteration 34: fitness improved to -24264.19 (swap key[l]<->key[s])
78 Iteration 35: fitness improved to -24247.82 (swap key[g]<->key[w])
79 Iteration 63: fitness improved to -24229.49 (swap key[g]<->key[u])
80 Iteration 123: fitness improved to -24227.94 (swap key[m]<->key[v])
81 Iteration 215: fitness improved to -24211.24 (swap key[h]<->key[l])
82 Iteration 275: fitness improved to -24180.50 (swap key[h]<->key[a])
83 Iteration 317: fitness improved to -24176.92 (swap key[r]<->key[a])
84 Iteration 479: fitness improved to -24155.21 (swap key[g]<->key[l])
85 Iteration 483: fitness improved to -23923.54 (swap key[i]<->key[r])
86 Iteration 625: fitness improved to -23921.11 (swap key[g]<->key[i])
87 No improvement for 1500 iterations; stopping.
88 Restart 1 (frequency-based) final fitness: -23921.11
89 Restart 2 (crib meh->the) final fitness: -22974.51
90 Restart 3 (frequency-based) final fitness: -24100.22
91 Restart 4 (crib meh->the) final fitness: -22977.02
92 Restart 5 (frequency-based) final fitness: -24020.09
93
94 =====
95 Exhaustive pair-swap refinement
96 =====
97 STEP: Try all 325 pair swaps in the key; keep any swap that improves n-gram
98 fitness.
99 JUSTIFICATION: This fixes remaining letter confusions (e.g. b/w, p/k) so the
100 decrypted text is fully coherent.

```

## 6.2 Manual phase: letter-swap corrections and final key

The manual letter-swap corrections were applied in this order, with the reason for each: (1) (1,22)  $b \leftrightarrow w$  (world/was); (2) (10,21)  $k \leftrightarrow v$  (marketing/governments); (3) (1,15)  $b \leftrightarrow p$  (product/described); (4) (9,23)  $j \leftrightarrow x$  (adjacently); (5) (16,23)  $q \leftrightarrow x$  (complex); (6) (16,25)  $q \leftrightarrow z$  (blitzkrieg, emphasized). The log for the manual phase and the final key and decryption is given below. The key and file paths shown here match those presented in Section 5.

```

1 =====
2 Manual letter-swap corrections for coherence
3 =====
4 STEP: Apply swaps ( $b \leftrightarrow w$ ), ( $k \leftrightarrow v$ ), ( $b \leftrightarrow p$ ), ( $j \leftrightarrow x$ ), ( $q \leftrightarrow x$ ), ( $q \leftrightarrow z$ ) for
5 full coherence.
6 JUSTIFICATION: Fixes world/was, marketing/governments, product/described,
7 adjacently, complex, blitzkrieg, emphasized.
8
9 (1,22)  $b \leftrightarrow w$ : corrects 'world' vs 'was' (e.g. 'end of world war i')
10 (10,21)  $k \leftrightarrow v$ : corrects 'marketing' vs 'governments' (e.g. 'military and
    government services')
11 (1,15)  $b \leftrightarrow p$ : corrects 'product' vs 'described' (e.g. 'finished product', ,
    described in detail')
12 (9,23)  $j \leftrightarrow x$ : corrects 'adjacently' (e.g. 'used ... adjacently')
13 (16,23)  $q \leftrightarrow x$ : corrects 'complex' (e.g. 'the most complex')
14 (16,25)  $q \leftrightarrow z$ : corrects 'blitzkrieg' and 'emphasized'

```

```

15 Applied swaps: (1,22)(10,21)(1,15)(9,23)(16,23)(16,25)
16 =====
17 Final substitution key and decryption
18 =====
19 STEP: Apply inverse substitution to ciphertext to obtain plaintext.
20 JUSTIFICATION: Each cipher letter is replaced by the corresponding plain
21 letter from the discovered key.
22
23
24 Key (plain -> cipher):
25 Plaintext a b c d e f g h i j k l m n o p q r s t u v w x y z
26 Ciphertext g t q d h b j e i v z r c w s p f x u m k o l y n a
27
28 Decrypted text saved to: .../data/task1/article_decrypted.txt
29 Key saved to: .../data/task1/substitution_key.txt

```

## 7 Observations and discussion

This section discusses observations from the decryption process and compares the approach with the SEED Lab exercise. The discussion draws on the methodology (Section 3), the results (Section 5), and the execution log (Section 6).

### 7.1 Observations

**Ciphertext format and encoding scope.** Inspection of the provided ciphertext and the decrypted output shows that **only the 26 letters (a–z) were encoded** by the substitution; all other characters are left unchanged by the decryption, consistent with the implementation. The ciphertext contains **no space characters, no digits, and no punctuation**. Numerals (e.g. years such as 1918 or 1923) do not appear in the ciphertext; they were likely omitted from the source before encryption or were not encoded—they are not present in the file and were not replaced by [] or any other symbol. The only non-letter characters present are the literal two-character sequence [] (in three positions) and newline characters; [] appears in the same positions in both ciphertext and decrypted text because the substitution is applied only to letters and those characters are preserved unchanged. Spaces were not encoded and are absent; the decrypted plaintext therefore has no word boundaries, and words run together (e.g. theenigmamachinewasinventedby...). **Case** is uniformly lowercase. These observations are consistent with a monoalphabetic substitution restricted to the 26 letters, with non-letters ([]) and newlines) preserved and spaces and numerals omitted from the ciphertext. They also justify the application of frequency analysis to a continuous letter stream, which is standard for this form of cryptanalysis [1, 2].

#### Heuristic and search behaviour.

- **Crib versus frequency-based start.** Using the most frequent cipher trigram (**meh**) as the yielded consistently higher final fitness than the purely frequency-based initial key, consistent with **the** being the dominant English trigram [1, 3].
- **Hill climbing.** Fitness improved in discrete steps (e.g. from approximately –24710 to approximately –22974 on the best run). Random pair swaps corrected letter confusions (e.g. similar-frequency letters). Termination after 1500 iterations without improvement avoided prolonged runs without gain.

- **Exhaustive refinement.** After hill climbing, exhaustive pair-swap refinement did not alter the key in the run shown in Section 6; the best key was already at a local maximum over the 325 possible single swaps. Manual swaps were still required for a small number of letter positions (e.g. b/w, k/v) that the n-gram model did not discriminate sufficiently [2].
- **Reference data.** The provided JSON files contained 26 letters, 42 bigrams, and 16 trigrams. Sparse trigram coverage is mitigated by the floor probability for unknown n-grams; richer reference data would be expected to improve discrimination.

## 7.2 Comparison with the SEED Lab exercise

The coursework asks whether the approach to discovering the encryption key was the same as in the lab exercise.

**No.** The SEED Lab on Secret-Key Encryption [5] focuses on the *use* of secret-key encryption: algorithms, modes of operation (e.g. ECB, CBC, CFB, OFB), padding, and initialisation vectors, using tools and programs to encrypt and decrypt when the *key is already known*. It does not address *cryptanalysis* of a substitution cipher or the *recovery* of an unknown key.

In Task 1, the key was *unknown* and the objective was to *recover* it from the ciphertext alone. The approach was therefore cryptanalytic: n-gram frequency analysis (monograms, bigrams, trigrams), a fitness function, and search (hill climbing with restarts and exhaustive refinement) were used to find the substitution key that yields English-like plaintext. The methods thus differ: the lab concerns the correct use of encryption and decryption with a given key, whereas Task 1 concerns the cryptanalysis of a classical cipher using language statistics and search. Both form part of the same coursework but serve different learning objectives—encryption versus cryptanalysis.

## 8 Conclusion

This report has presented the approach, implementation, and results for breaking the substitution cipher in Task 1. In summary, the cipher was broken by: (i) loading letter, bigram, and trigram frequencies from the provided JSON data; (ii) constructing an initial key from cipher letter counts matched to English letter frequency order, and optionally from a crib (`meh` → `the`); (iii) maximising an n-gram log-probability fitness via hill climbing with random pair swaps and multiple restarts; (iv) exhaustive pair-swap refinement; and (v) a small number of manual letter swaps for coherence. The encryption key was recovered and the ciphertext decrypted to readable plaintext about the Enigma machine. It is concluded that monoalphabetic substitution is vulnerable to statistical attack when sufficient ciphertext and language-specific n-gram data are available [2], and that the combination of frequency-based initial keys, cribs, and local search (hill climbing with restarts) constitutes an effective and well-established cryptanalytic approach [1, 2].

## References

- [1] J. Kun, “Cryptanalysis with N-grams.” Math & Programming, 2012. Accessed for this report. Discusses substitution cipher representation, trigram and bigram scoring, steepest ascent, and bigram-guided neighbour generation.

- [2] Practical Cryptography, “Cryptanalysis of the simple substitution cipher.” Practical Cryptography, n.d. Describes hill climbing with quadgram fitness, random key and swap steps, multiple restarts, unicity distance, and limitations.
- [3] P. Norvig, “Letter n-gram counts.” count\_2l.txt, count\_3l.txt, n.d. Widely used for bi-gram/trigram statistics; referenced in Kun (2012).
- [4] Code Review Stack Exchange, “Decrypting a substitution cipher using n-gram frequency analysis.” Stack Exchange, n.d. Discussion and code for n-gram-based decryption.
- [5] SEED Labs, “Secret-key encryption.” SEED Labs 2.0 – Crypto Lab, 2020. Lab on encryption algorithms, modes, padding, and IV; uses tools and programs to encrypt/decrypt messages.